

# NEW YORK CITY TAXI TRIP DURATION

TIANJIAO WANG

ABSTRACT. This project is to build a model that predicts the total ride duration of taxi trips in New York City. The primary dataset is one released by the NYC Taxi and Limousine Commission, which includes pickup time, geo-coordinates, number of passengers, and several other variables.

## CONTENTS

1. Introduction	2
2. Data loading and overview	2
3. Data cleaning	3
4. Features engineering	4
5. Model selection	5
6. Hyperparameters tuning	6
7. Training and predictions	6

---

*Date:* (None).

*2020 Mathematics Subject Classification.* Artificial Intelligence.

*Key words and phrases.* New York, Taxi, Predict.

## 1. INTRODUCTION

The project will build a model that predicts the total ride duration of taxi trips in New York City. The primary dataset is one released by the NYC Taxi and Limousine Commission, which includes pickup time, geo-coordinates, number of passengers, and several other variables. Accordingly, this project problem is taxi trips duration, which is a outlier detection. There are 6 steps including data loading and overview, data cleaning, features engineering, model selection, hyperparameters tuning, and training and predictions.

- Data loading and overview
- Data cleaning
- Features engineering
- Model selection
- Hyperparameters tuning
- Training and predictions

## 2. DATA LOADING AND OVERVIEW

At first, I quickly look at the first 5 lines of a dataset to understand the structure, format, and content of the data . Then I take a overview of the type and amount and other information of df and test data.

Colonne	Description
<b>id</b>	a unique identifier for each trip
<b>vendor_id</b>	a code indicating the provider associated with the trip record
<b>pickup_datetime</b>	date and time when the meter was engaged
<b>dropoff_datetime</b>	date and time when the meter was disengaged
<b>passenger_count</b>	the number of passengers in the vehicle (driver entered value)
<b>pickup_longitude</b>	the longitude where the meter was engaged
<b>pickup_latitude</b>	the latitude where the meter was engaged
<b>dropoff_longitude</b>	the longitude where the meter was disengaged
<b>dropoff_latitude</b>	the latitude where the meter was disengaged
<b>store_and_fwd_flag</b>	This flag indicates whether the trip record was held in vehicle memory before sending to the vendor because the vehicle did not have a connection to the server (Y=store and forward; N=not a store and forward trip)
<b>trip_duration</b>	duration of the trip in seconds

FIGURE 1. overview of the data



### 3. DATA CLEANING

I do the data learning to check the duplicated and missing values and deal with outliers.

```

id                0
vendor_id         0
pickup_datetime   0
dropoff_datetime  0
passenger_count   0
pickup_longitude  0
pickup_latitude   0
dropoff_longitude  0
dropoff_latitude  0
store_and_fwd_flag 0
trip_duration     0
dtype: int64

```

FIGURE 2. No duplicate and missing data

- Visualize outliers

There are outliers. I can't find a proper interpretation and it will probably damage our model, so I choose to get rid of them.




FIGURE 3. boxplot for trip-duration




FIGURE 4. pickup positions




FIGURE 5. dropoff positions

In this step, I only keep trips that lasted less than 5900 seconds, and only keep trips with passengers, and remove position outliers (pickup-longitude more than -100, pickup-latitude less than 50; dropoff-longitude less than -70 and dropoff-longitude more than -80, dropoff-latitude less than 50).

#### 4. FEATURES ENGINEERING

At first, I visualize the distribution of trip-duration values.




FIGURE 6. the distribution of trip-duration values

The distribution is right-skewed so we can consider a log-transformation of trip-duration column. When dealing with a right-skewed distribution, considering a log-transformation of the data can be a useful technique. Right-skewed distributions are characterized by a long tail on the right side, where some data points have very high values, leading to a positively skewed distribution.

Logarithmic transformation, specifically taking the natural logarithm (base e) or common logarithm (base 10) of the data, is a common approach to mitigate the skewness in the distribution.

Then I deal with categorical features using One-hot encoding binary categorical features. One-hot encoding is a common technique used in data preprocessing, particularly in machine learning, for handling categorical data. It's particularly useful when the categorical data is nominal (i.e., there is no intrinsic order to the categories).

Then I deal with dates, and creat distance and speed. Finally, I remove the outliers. Then we test the correlations between variables and made a figure.

🔥 (None)-(None) ((None))




FIGURE 7. log of trip-duration




FIGURE 8. boxplot for distance




FIGURE 9. boxplot for speed




FIGURE 10. correlations between variables

## 5. MODEL SELECTION

In this part, we split the labeled data frame into two sets: features and target to train, and then test the models. Then, for this specific problematic, we'll

🔥 (None)-(None) ((None))

5

Committed by: (None)

measure the error using the RMSE(Root Mean Squared Error). We tried Gradient-Boosting, RandomForest, and LightGBM. LightGBM is blazingly fast compared to RandomForest and classic GradientBoosting, while fitting better. It is our clear winner. Then we did cross-validation. The result shows that our LightGBM model is stable.

## 6. HYPERPARAMETERS TUNING

We did hyperparameters tuning using RandomizedSearchCV, and Test the following parameters: 'learning-rate', 'max-depth', 'num-leaves', 'objective': 'regression', 'feature-fraction', 'bagging-fraction', 'max-bin'.

## 7. TRAINING AND PREDICTIONS

We trained on all labeled data using the best parameters in hyperparameters tuning, and trained on all labeled data using the best parameters (sklearn API version), and trained on all labeled data using the best parameters. Then we make predictions on test data frame, create a data frame designed a submission on Kaggle. Lastly, create a csv out of the submission data frame.

	id	trip_duration
0	id3004672	716.070826
1	id3505355	672.125770
2	id1217141	455.368356
3	id2150126	938.637832
4	id1598245	354.432595

FIGURE 11. predict-result

(A. 1) BUSINESS SCHOOL,, BEIJING TECHNOLOGY AND BUSINESS UNIVERSITY, BEIJING 100048, CHINA

Email address, A. 1: wangtianjiao@gmail.com